

Project 4 Loan Prediction

Import necessary libraries

```
In [3]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report
```

Step 1: Load the dataset

```
In [6]: # Step 1: Load the dataset
data_path = r"C:\Users\Hello\Desktop\Loan_Predication.csv"
data = pd.read_csv(data_path)
```

Step 2: Explore the data (EDA)

```
In [11]: # Step 2: Explore the data (EDA)
print("Dataset Overview:")
print(data.info())
print("\nMissing Values:")
print(data.isnull().sum())
print("\nSummary Statistics:")
print(data.describe())
```

Dataset Overview:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 614 entries, 0 to 613

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	Loan_ID	614 non-null	object
1	Gender	601 non-null	object
2	Married	611 non-null	object
3	Dependents	599 non-null	object
4	Education	614 non-null	object
5	Self_Employed	582 non-null	object
6	ApplicantIncome	614 non-null	int64
7	CoapplicantIncome	614 non-null	float64
8	LoanAmount	592 non-null	float64
9	Loan_Amount_Term	600 non-null	float64
10	Credit_History	564 non-null	float64
11	Property_Area	614 non-null	object
12	Loan_Status	614 non-null	object

dtypes: float64(4), int64(1), object(8)

memory usage: 62.5+ KB

None

Missing Values:

Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0

dtype: int64


Summary Statistics:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
count	614.000000	614.000000	592.000000	600.000000

mean	5403.459283	1621.245798	146.412162	342.000000
std	6109.041673	2926.248369	85.587325	65.12041
min	150.000000	0.000000	9.000000	12.000000
25%	2877.500000	0.000000	100.000000	360.000000
50%	3812.500000	1188.500000	128.000000	360.000000
75%	5795.000000	2297.250000	168.000000	360.000000
max	81000.000000	41667.000000	700.000000	480.000000

Credit_History	
count	564.000000
mean	0.842199
std	0.364878
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

```
In [13]: # Visualizations
sns.countplot(x='Loan_Status', data=data)
plt.title("Loan Status Distribution")
plt.show()
```



Step 3: Handle missing values

```
In [16]: # Step 3: Handle missing values
imputer = SimpleImputer(strategy='most_frequent')
data_imputed = pd.DataFrame(imputer.fit_transform(data), columns=data.columns)
```

Step 4: Encode categorical variables

```
In [19]: # Step 4: Encode categorical variables
categorical_cols = data_imputed.select_dtypes(include=['object']).columns
label_encoders = {}
for col in categorical_cols:
```


```
label_encoders[col] = LabelEncoder()  
data_imputed[col] = label_encoders[col].fit_transform(data_imputed[col])
```

Step 5: Feature engineering

```
In [22]: # Step 5: Feature engineering  
data_imputed['Income_to_Loan_Ratio'] = data_imputed['ApplicantIncome'] / (data_imputed['LoanAmount'] + 1)  
data_imputed['Family_Size'] = data_imputed['Dependents'].replace('3+', 3).astype(int) + 1
```

Step 6: Feature selection

```
In [25]: # Step 6: Feature selection  
corr_matrix = data_imputed.corr()  
plt.figure(figsize=(10, 8))  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')  
plt.title("Feature Correlation Matrix")  
plt.show()
```



C

Loi

Incon

Step 7: Split data

```
In [28]: # Step 7: Split data
X = data_imputed.drop('Loan_Status', axis=1)
y = data_imputed['Loan_Status']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Step 8: Build a baseline model (Logistic Regression)

```
In [31]: # Step 8: Build a baseline model (Logistic Regression)
logistic_model = LogisticRegression(max_iter=1000)
logistic_model.fit(X_train, y_train)
y_pred_baseline = logistic_model.predict(X_test)
```

C:\Users\Hello\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Step 9: Model evaluation (Logistic Regression)

```
In [34]: # Step 9: Model evaluation (Logistic Regression)
print("Logistic Regression Performance:")
print("Accuracy:", accuracy_score(y_test, y_pred_baseline))
print("Precision:", precision_score(y_test, y_pred_baseline))
print("Recall:", recall_score(y_test, y_pred_baseline))
print("F1 Score:", f1_score(y_test, y_pred_baseline))
```

Logistic Regression Performance:
Accuracy: 0.7837837837837838
Precision: 0.7564102564102564
Recall: 0.9833333333333333
F1 Score: 0.855072463768116

Step 10: Model tuning (Random Forest)

```
In [37]: # Step 10: Model tuning (Random Forest)
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)
```

```
In [39]: # Model evaluation (Random Forest)
print("\nRandom Forest Performance:")
print("Accuracy:", accuracy_score(y_test, y_pred_rf))
print("Precision:", precision_score(y_test, y_pred_rf))
print("Recall:", recall_score(y_test, y_pred_rf))
print("F1 Score:", f1_score(y_test, y_pred_rf))
```


Random Forest Performance:
Accuracy: 0.7891891891891892
Precision: 0.7682119205298014
Recall: 0.9666666666666667
F1 Score: 0.8560885608856088

```
In [41]: # Compare results
print("\nClassification Report (Random Forest):")
print(classification_report(y_test, y_pred_rf))
print("Confusion Matrix:")
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix (Random Forest)")
plt.show()
```


Classification Report (Random Forest):

	precision	recall	f1-score	support
0	0.88	0.46	0.61	65
1	0.77	0.97	0.86	120
accuracy			0.79	185
macro avg	0.83	0.71	0.73	185
weighted avg	0.81	0.79	0.77	185

Confusion Matrix:



Step 11: Analyze results

Step 11: Analyze results

```
print("Logistic Regression vs Random Forest:") print(f"Logistic Regression F1: {f1_score(y_test, y_pred_baseline)}") print(f"Random  
Forest F1: {f1_score(y_test, y_pred_rf)}")
```

In []: